

T0 · toolchain: Triton (GPU)
and a pinned triton-cpu
just triton-smoke on both
backends



T1 · committed profile
gprof_table.sh on ww3_ts1,
NK/NTH 36/24 and 32/36



T2 · LabECO: which
SDSCUM branch?



T3 · W3SDS4 fixture from
real sea points
+ a NumPy reference



no

NumPy passes L1?



yes

T4 · Triton G: literal
translation,
no improvements



no

L1 with a
justified tolerance?



yes

T5 · Triton C: same source
on triton-cpu,
every line that has to differ
is logged



builds and
passes L1?

yes



T6 · benchmark:
kernel vs kernel, projected
model gain



no: recorded as a result



T8 · report: is Triton worth
continuing?
a negative answer counts



T7 · next routine by wall
time:
W3SNL1, W3SIN4, then one
W3SRCE kernel