Wave models on GPUs

C++/Kokkos and Rust ports of WAVEWATCH III that must give the Fortran's bits

Matheus Hoffmann ORCID 0009-0009-1056-7661
Escola Politécnica, Universidade Federal do Rio de Janeiro, Brazil

Pedro Veras Guimarães ORCID 0000-0002-1686-4722
Co-advisor of the first author. LabECO, Departamento de Engenharia Mecânica, Universidade Federal de Santa Catarina, Brazil

Pedro Cruz ORCID 0000-0002-2019-133X
Advisor of the first author. Grupo de Teleinformática e Automação (GTA), DEL, Escola Politécnica, Universidade Federal do Rio de Janeiro, Brazil

11 November 2026 · v0 (preprint, not peer reviewed) · rev. 565e849+853

Abstract. We report two ports of the WAVEWATCH III (WW3) spectral wave model. Coding agents wrote both ports, and both ports have one acceptance rule: the port must give the same bits as the Fortran. A tolerance is a diagnostic, never the verdict. *ww3-gpu* keeps WW3, and the Fortran calls C++ written with Kokkos, behind one run-time flag for each routine. Its reference is the Fortran path of the same executable. *chaceon* replaces WW3 with Rust, compiled for the CPU, for SPIR-V (rust-gpu) and for PTX (Rust-CUDA). No Fortran runs in its process, and its reference is a WW3 that is compiled separately. ww3-gpu ports the discrete interaction approximation (DIA) of the four-wave transfer. On the CPU, it is byte-identical on all 100 serial and 25 OpenMP entries of the WW3 regression matrix that select it. The DIA takes 48.4 % of single-thread self time on a case of 9 520 points. chaceon runs `ww3_grid`, `ww3_strt`, `ww3_prep` and `ww3_shel` without Fortran, and 96 output files are byte-identical to the reference. Neither port has run the whole model on a real GPU, and this note claims no speed-up. We compare how each project defines and checks its evidence.

Keywords: WAVEWATCH III; bit-for-bit reproducibility; Kokkos; Rust; GPU porting; coding agents.

Preface

This project exists to support Ondas & Ventos (Universidade Federal de Santa Catarina 2026), an operational wave forecast system at the Federal University of Santa Catarina (UFSC). LabECO at UFSC runs WW3 operationally for ReNOMO, the Brazilian national network for ocean observation and monitoring (LabECO/UFSC 2025). ReNOMO started in 2023 with the CNPq/MCTI/Finep call 062/2022. More than 100 researchers from 31 Brazilian institutions take part in it (LabECO/UFSC 2025).

The system must run the wave model at each forecast cycle, and the run time of WW3 limits the cycle. The system also plans to force the wave model with more than one wind model: the IFS of ECMWF, the GFS of NCEP and Google's WeatherNext ⚠. Each wind model adds one full WW3 run, so an ensemble of wind models multiplies the cost of the cycle. A WW3 that runs faster on GPUs and gives the same bits makes these runs affordable without a change to the forecast. ww3-gpu and chaceon work toward that WW3, and this note claims no speed-up.

0 Context: an ocean model ported by a language model

Koldunov et al. (Koldunov et al. 2026) report the port of a whole ocean model by a large language model. The model is FESOM2, an ocean and sea-ice model, and the authors are at the Alfred Wegener Institute. Claude Code, with Claude Opus 4.7, translated it from Fortran to C and then from C to C++/Kokkos. The people set the strategy, reviewed the plans and the diffs, and found physics and validation errors. The agent read, translated, built the test harnesses, debugged, profiled and wrote the documentation.

The C stage is the centre of their method. The clean C code is the reference, and each Kokkos kernel keeps its C original as a "twin". The acceptance rule changes at each step. The C model is compared with the Fortran by statistics: over five years, the RMS difference of the mean is 0.006 °C in sea-surface temperature. Only the Kokkos Serial model is byte-identical, and only to the C model. The authors write that bit-for-bit agreement between Fortran and C "across compilers and arithmetic orderings is neither attainable nor necessary". On the GPU, they accept differences up to a recorded ceiling for each field.

The successor of WW3 is also in progress. NCEP Office Note 525 (Tolman 2025) chose a bottom-up rewrite, with a C++ core first and a Rust core later ⚠. The same report demotes Fortran to a solver-only language ⚠.

The WAVEWATCH IV repository (NOAA Environmental Modeling Center 2026) is pre-alpha: on 2026-11-11 its `develop` branch had 40 commits, from 2025-11-17 to 2026-10-16 (`875f49a`). Its 97 source files are C++ (66 `.cpp`, 31 `.h`), and no physics kernel of WW3 is in it yet.

Thus the core kernels of the wave model stay in Fortran for some years ⚠. This is the first reason why the first author will continue to maintain ww3-gpu and chaceon. They make the Fortran kernels of WW3 run on GPUs now, without a change to their bits.

This note starts from that result and asks a narrower question. Can an agent port a spectral wave model, and give the bits of the Fortran itself? The two ports below have no C stage, and their reference is the Fortran. They also use two practices that FESOM2 reports: a strictly literal translation, and a `file:line` with the literal value for each statement about the Fortran.

1 Why bit for bit

WW3 integrates the action balance equation

$$\frac{\partial N}{\partial t} + \nabla \cdot (\dot{\mathbf{c}}\, N) = \frac{S}{\sigma},$$

where the source term S is the sum of wind input, dissipation, nonlinear transfer and bottom terms. A port that changes the order of a floating-point sum changes N. A forecast system that was validated against observations must then be validated again. Thus both projects keep the arithmetic of the Fortran. They do not reorder a reduction, they do not use a fused multiply-add (`-ffp-contract=off` on both sides), and they do not widen a tolerance. A difference is a defect of the port until two runs of the Fortran itself give different bytes in the same file.

2 Two shapes for one rule

Table 1: The two ports. Sources: chaceon `docs/WW3_GPU_PARITY_202611.md` (2026-11-09) and `README.md` at `0c098fe` (Hoffmann 2026a); ww3-gpu `kokkos/PORT_STATUS.md` at `58167d7` (Hoffmann 2026c).

	ww3-gpu (C++/Kokkos)	chaceon (Rust)
Shape	Inside WW3: the Fortran calls C++ through a `bind(C)` shim, one flag for each routine	Replaces WW3: no Fortran in a chaceon process
Reference	NOAA-EMC/WW3 `03adda91`, through a fork that adds 5 687 lines and removes none	MetOcean WW3 7.14, `92e2cf66`
CPU parity against	the Fortran path of the same executable	a reference that is compiled separately
Device targets	Kokkos 5.2.2, CUDA backend (`sm_89`); OpenMP on the CPU	rust-gpu (SPIR-V, Vulkan) and Rust-CUDA (PTX)
Spectrum resident on the device	yes, on the OpenMP backend	yes, in SPIR-V on CPU Vulkan (lavapipe)
Size of the port	62 257 lines of C++, 355 files (2026-10-14)	29 229 lines of Rust port, 5 936 lines of tools
Formal methods	SMT queries, exhaustive sweeps	one Lean 4 proof about `f32` division
Licence	MIT, LGPL-3.0 kernels	MPL-2.0

Table 2: The two repositories. Source: the GitHub REST API (`repos/h0ffmann/<repo>` for `created_at` and `pushed_at`, `commits?sha=main` for the commits), read on 2026-11-11 at 03:35 UTC. The first commit of ww3-gpu is older than the repository because it was made before the first push.

Metric	ww3-gpu	chaceon
Repository created	2026-09-12 01:31 UTC	2026-10-23 03:58 UTC
First commit on `main`	2026-09-11 13:47 UTC, `b43a6d3`	2026-10-23 03:58 UTC, `3bad31a`
Last commit on `main`	2026-11-11 03:16 UTC, `efc0211`	2026-11-09 08:53 UTC, `0c098fe`
Last push, any branch	2026-11-11 03:31 UTC	2026-11-09 08:53 UTC
Commits on `main`	800	110

The Kokkos route keeps every WW3 program and replaces one routine at a time. Thus one regtest can isolate the effect of one routine. The Rust route must rebuild the programs, the file formats and the time loop before it can compare a whole run. After that, no Fortran remains that can share state with the port.

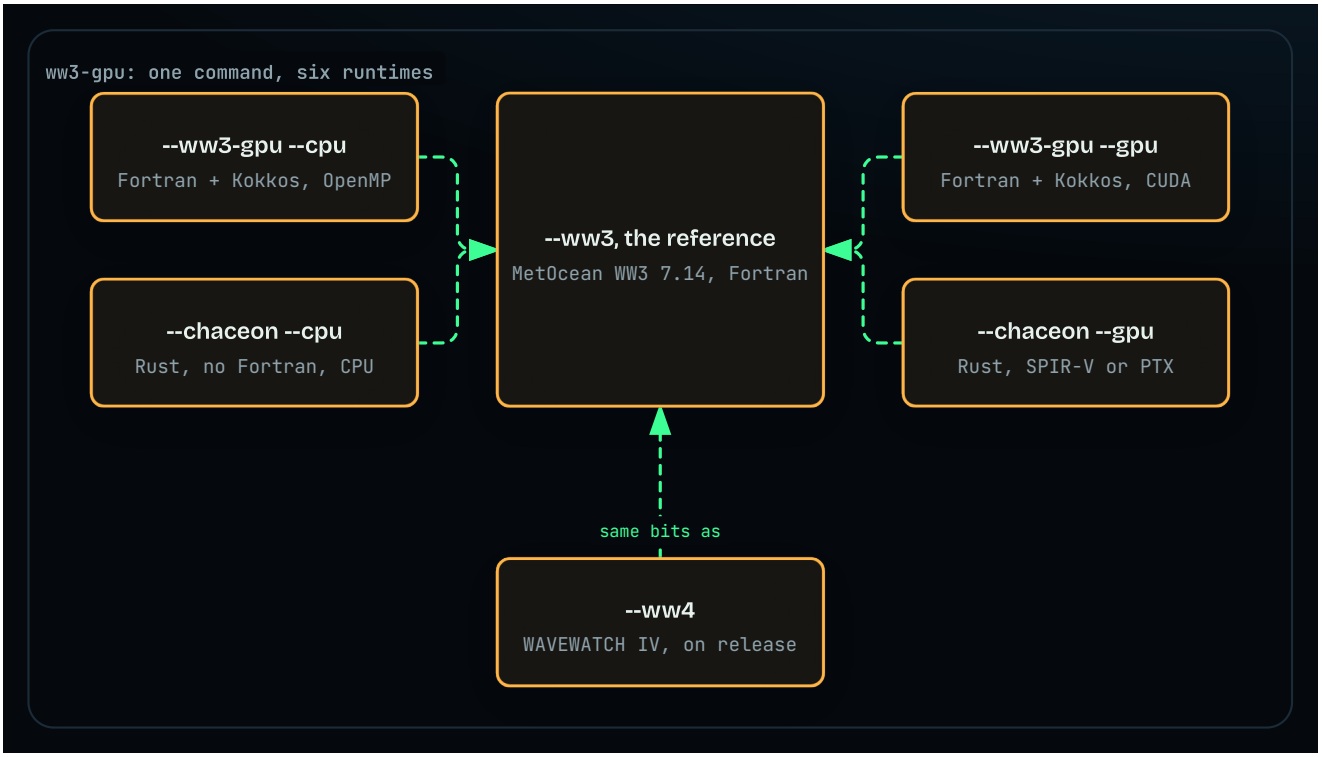


Figure 1: The six runtimes of ww3-gpu. A dashed arrow points to the run that must give the same bits: all five point to `--ww3`, the reference. Drawn with Archify from `docs/archify/ww3-gpu-runtimes.json` (Hoffmann 2026c).

The CPU results of the whole ww3-gpu runtime, with the resident spectrum, and their limits are in its NOAA progress record (Hoffmann 2026e). The DIA port is also published alone as ww3-dia-gpu-v2 (Hoffmann 2026d). That repository compares the port with an earlier CUDA kernel of the same routine.

Figure 1 shows the six runtimes of one command, `scripts/ww3-gpu.sh`. `--ww3` runs the Fortran of WW3 and is the reference: each other runtime must give its bits. `--ww3-gpu` builds the fork with the KOKKOS switch and links the Kokkos port. With `--cpu` it runs on OpenMP, and with `--gpu` it runs on CUDA. In that executable, one flag for each routine selects the Fortran or the Kokkos code. `--chaceon`, or `--cceon`, runs WW3 rewritten in Rust, with no Fortran in the process, on the CPU or on a GPU. `--ww4` is for WAVEWATCH IV when it is released. Today it builds the development code from a fork, and it stops with exit code 5 when the case needs a part that is not ported. Two parts of this design are not in the script yet ⚠. The script has no `--chaceon` flag. And `--ww3` builds NOAA WW3 `03adda91`, while chaceon compares with MetOcean WW3 7.14, compiled separately. The design gives the six runtimes one reference, MetOcean WW3 7.14.

3 Evidence ladders

Each project gives names to its steps. A step in one ledger is not the same evidence as a step in the other ledger.

Table 3: Evidence steps. ww3-gpu counts come from `scripts/port_score.py --write` (profile `bench_small`, 120 × 80 points, 12 h, one thread, gprof, 2026-10-14); chaceon counts come from `just port-summary`.

Step	ww3-gpu	chaceon
Code exists	*translated*: 316 routines, 26.2 % of Fortran lines	*translated*: Rust exists, not fully reached
A run calls it	*ported*: 72 routines, 99.3 % of profiled run time	*routine*: every captured call is bit-identical on the host and in SPIR-V
Whole runs identical, CPU	*regtested*: 14 routines, 48.6 % of run time	*host*: whole runs are `cmp`-identical; 10 080 lines
Identical through a GPU API	none	*spirv*: the same files from SPIR-V on CPU Vulkan; 5 189 lines
Identical on a real GPU	*gpu*: none	*cuda*: none

4 What each project found

ww3-gpu: the reference is the optimised code. Byte identity for the DIA needed one fact that the source text does not give. gfortran expands `x**11` as one chain of multiplications at `-O0` and as a different chain at `-O3`. The two chains disagree on 643 427 of 1 000 000 single-precision inputs. Before this was found, the research log explained the remaining differences as round-off that WW3 amplifies. That explanation was wrong. Thus the reference of the port is the code that the WW3 build produces, read as GIMPLE or disassembly, not the Fortran text. The port has no unit tests of its own. Only the WW3 regression tests, compared by `test.comp`, judge it.

chaceon: a run with no Fortran in it. The Rust programs `ww3_grid`, `ww3_strt`, `ww3_prep` and `ww3_shel` run in sequence from the namelists. They write `mod_def.ww3`, the forcing, the gridded and point outputs, the track file and both restarts. All 96 files are byte-identical to the reference. The boundaries give 214 identical comparisons over 34 cases, and the track post-processors give 74. `ww3_shel --threads N` and an MPI build give the same bytes. The files stay identical when `W3SRCE`, `W3XYP3` and `W3KTP3` are compiled to SPIR-V and run on a CPU Vulkan driver. Enumeration cannot reach the 2^{64} pairs of `f32` operands. There, a Lean proof states the division fact that the port uses. The proof says nothing about whether Rust equals Fortran, because Fortran has no usable formal semantics.

Why ww3-gpu has no proof yet, and what comes next. In this first phase, ww3-gpu adopted no proof assistant (Hoffmann 2026f). There are four reasons:

1. No proof assistant reads Fortran or C++. A proof is about a model that a person writes. It says nothing about the gfortran and nvcc binaries unless somebody shows that the model agrees with them, and that agreement is the open question.
2. Bit equality is a property of two binaries, built with given flags. It is not a property of two source texts.
3. For straight-line IEEE-754 code, equality on every input is decidable. SMT queries in the `QF_FP` theory and sweeps over all 2^{32} `float32` inputs answer it in seconds to minutes. An interactive proof takes months for each routine.
4. Agents must write each proof, and a person must review it. This is a second artefact for each routine.

Without a proof, the DIA port has the strongest evidence short of one. It is byte-identical on all 125 entries of the WW3 regression matrix that select it (100 serial, 25 OpenMP). Section 1 of the kernel was also swept over all 2^{32} `float32` inputs, with 0 differences (Hoffmann 2026f). Thus the DIA port is validated by the tests of the model that it replaces, on every entry that calls it. This validation is not a formal proof, and the note does not call it one.

A first formal layer is already next to the ports. The Lean 4 package `ocean-waves` (Hoffmann 2026b) states the discretised action balance equation and proves these properties, with no `sorry` and no `axiom`:

- $\sigma^2(k)$ increases strictly with k;
- $c/2 \le c_g \le c$;
- the depth CFL condition keeps every Courant number in $[0, 1]$;
- upwind propagation and a DIA-shaped quadruplet exchange conserve action;
- the semi-implicit source step keeps the action positive for every time step.

The package uses exact `Rat` arithmetic on a reduced model: one space dimension, a fixed depth and a Padé model of tanh. Thus it also says nothing about bits. Agda is the next step. Agda has no 32-bit float type, so it cannot prove a fact about bits. Its role is a total, executable reference semantics of each source term. Each table index will have a `Fin` type, so every read is in bounds for every grid, not only for one grid. The sweeps, the solver queries and the WW3 regression tests continue to decide the bits.

5 Challenges: the regression tests are large

WW3's regression tests are the only judge of the Kokkos port, and they are large. Table Table 4 gives their size at the fork commit `80644ba9`, counted on 2026-10-18.

Table 4: Size of the WW3 regression matrix. Source: ww3-gpu `regtests/README.md` at `58167d7`, counts from `just ww3-gpu-matrix-list`. The count of all entries is the number of lines that `matrix.base` writes, and a machine section can disable some of them.

Set of `matrix.base` entries	Entries	Share
All entries	758, in 60 tests	100 %

Set of `matrix.base` entries	Entries	Share
Entries with the DIA (`NL1`)	310	41 %
Entries with ST4 physics (`ST4`)	136	18 %
Entries that the resident preset can select	7 (all serial)	0.9 %
Selected regtests of ww3-gpu	5, and the `ww3_ts1` pair	0.7 %

One run of the 310 `NL1` entries takes days. The selected set takes one to three hours on one workstation (Intel i9-14900, 32 hardware threads); this time is an estimate, not a measurement ⚠. Thus a full matrix run cannot follow each commit. ww3-gpu runs the selected set as a guardrail, and it postpones the full matrix to a validation phase. Until then, a static argument supports bit parity: the `-O3` GIMPLE or the disassembly of both compilers.

6 What is not claimed

1. Neither project has run the whole model on a real GPU. The PTX of chaceon has never run. ww3-gpu has run one DIA benchmark kernel on an RTX 4090 (Hoffmann 2026d).
2. This note claims no speed-up. The rules of ww3-gpu mark its recorded timings (Hoffmann 2026e) "regtest pending", and chaceon has no timings.
3. On a GPU, bit identity also needs the `expf` and `logf` of the device to return the bits of glibc. ww3-gpu sends its kernels to device copies of the glibc 2.44 functions. These copies equal the host functions on every input that was checked. No device run has compared them yet. Until one does, a fixed tolerance judges a GPU run.
4. The two references are different lines of WW3 (NOAA-EMC and MetOcean). Thus the comparison of the two ports is by routine name, not by source text. Both ledgers contain 85 routines.

7 Next steps

The authors plan four steps:

1. Continue the development of ww3-gpu and chaceon on an NVIDIA RTX 4090.
2. Run the selected regression tests on an NVIDIA H100 that UFSC owns.
3. Pin every runtime to WW3 7.14.
4. Publish version 1.0.0 of ww3-gpu and of chaceon before February 2027.

These steps are plans, not results ⚠. The first author supervises each GPU run, as the rules of ww3-gpu require.

8 Use of AI

Claude coding agents (Anthropic) produced the C++/Kokkos and Rust code, the shims, the comparison scripts, the investigations and drafts of this text. They worked under the direction of the first author, in the two repositories. Pedro Cruz (GTA, UFRJ) is the advisor and Pedro Veras Guimarães (LabECO, UFSC) is the co-advisor of the first author.

Most of the commits came from Claude Opus. Claude Sonnet also ported a large amount of code without gross errors, in the experience of the first author ⚠. One example is the split of `W3SDS4` into parts that run one sea point per vector lane (ww3-gpu commit `9cef132`, 841 lines added). Its output files were byte-identical to the Fortran of the same executable on the `smoke` and `full` cases.

Each repository publishes the rules, skills and hooks of its agents in its `AGENTS.md` file and its `.claude/` directory. Every number in this note comes from a file of a repository at the commit that the note names. Nothing was run for this note. Each claim was checked on 2026-11-11 against the source and the commit that it names; a ⚠ marks a claim that was not checked against a primary source. The physics of WW3 is used without change; WAVEWATCH III is a registered trademark of NOAA.

Disclaimer. Today, both repositories are set up for Claude Code. The authors will continue to work so that any coding agent can use ww3-gpu and chaceon. This includes agents that run local models with Ollama. This work is not complete, and this note does not claim results for other agents or for local models.

References

Hoffmann, Matheus. 2026a. “chaceon: A Rust Port of WAVEWATCH III.” https://github.com/h0ffmann/chaceon.

———. 2026b. “Ocean Waves: The Physics of a Spectral Wave Model in Lean 4.” https://github.com/h0ffmann/proof/tree/b863c15/ocean-waves.

———. 2026c. “WW3 GPU Lab: Agentic Coding and Agentic Research on a Bit-for-Bit GPU Port of WAVEWATCH III.” https://doi.org/10.5281/zenodo.23221351.

———. 2026d. “ww3-dia-gpu-v2: The Whole W3SNL1, Inside WAVEWATCH III, Bit for Bit.” https://github.com/h0ffmann/ww3-dia-gpu-v2/tree/32aa2ef.

———. 2026e. “A Bit-Identical Kokkos Runtime for WAVEWATCH III: CPU Results and Five Speed-Ups.” ww3-gpu, `docs/NOAA_PROGRESS_RECORD_202610.md`, commit 58167d7. https://github.com/h0ffmann/ww3-gpu/blob/58167d7/docs/NOAA_PROGRESS_RECORD_202610.md.

———. 2026f. “ADR-0001: No Proof Language for the GPU Port; SMT Queries and Exhaustive Sweeps Instead.” ww3-gpu, `docs/ADRs/ADR-0001-proof-language.md`, commit 58167d7. https://github.com/h0ffmann/ww3-gpu/blob/58167d7/docs/ADRs/ADR-0001-proof-language.md.

Koldunov, Nikolay V., Suvarchal K. Cheedela, Sergey Danilov, Dmitry Sidorenko, Sebastian Beyer, and Thomas Jung. 2026. “An Ocean Model Ported by a Large Language Model: Experience and Lessons from FESOM2 (Fortran to C to C++/Kokkos).” arXiv:2606.11356 [physics.ao-ph]. https://doi.org/10.48550/arXiv.2606.11356.

LabECO/UFSC. 2025. “ReNOMO: Rede Nacional de Observação e Monitoramento Oceânico.” https://labeco.ufsc.br/renomo/.

NOAA Environmental Modeling Center. 2026. “WAVEWATCH IV.” https://github.com/NOAA-EMC/WW4.

Tolman, Hendrik L. 2025. “The WAVEWATCH III Software Modernization Project: Phase I Report.” Office Note 525. College Park, MD: NOAA/NWS/NCEP. https://doi.org/10.25923/h7j3-1h25.

Universidade Federal de Santa Catarina. 2026. “Ondas & Ventos.” https://ondaseventos.ufsc.br/.